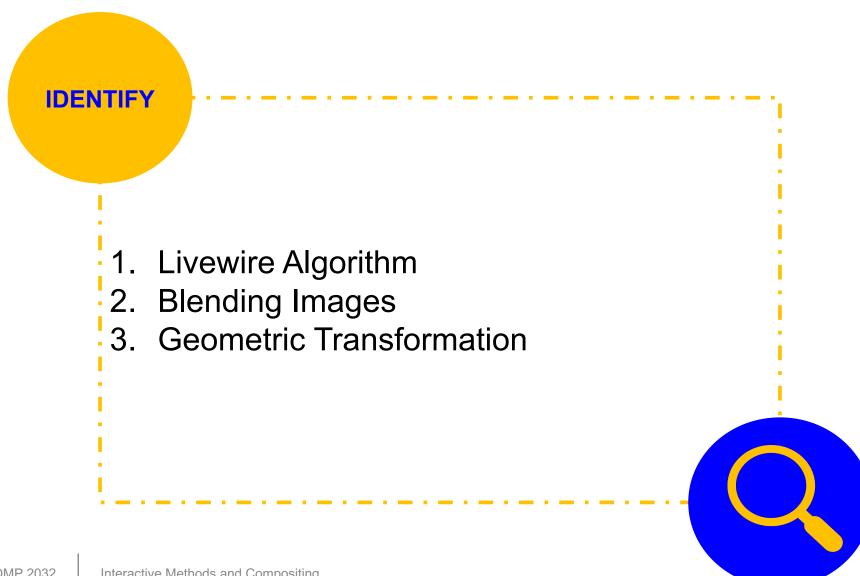


Introduction to Image Processing

Lecture 11
Interactive Methods and Compositing



Learning Outcomes





LiveWire Algorithm



Compositing



© NASA



Compositing

Extract sprites using e.g. intelligent scissors in Photoshop











Blend them into the composite (in the right order)



Intelligent Scissors

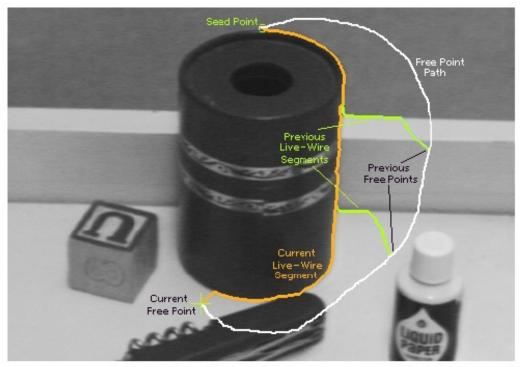


Figure 2: Image demonstrating how the live-wire segment adapts and snaps to an object boundary as the free point moves (via cursor movement). The path of the free point is shown in white. Live-wire segments from previous free point positions $(t_0, t_1, and t_2)$ are shown in green.

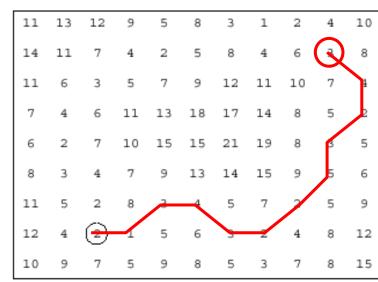
E. N. Mortensen and W. A. Barrett, Intelligent Scissors for Image Composition, in *ACM Computer Graphics (SIGGRAPH `95)*, 1995



Intelligent Scissors

We need a path from seed to mouse that follows the object boundary as closely as possible

- Define a path that stays as close as possible to edges
- Quantify the "cost" of including each pixel in the path
- Edge pixels have low cost
- Find lowest cost path from seed to cursor



mouse

seed



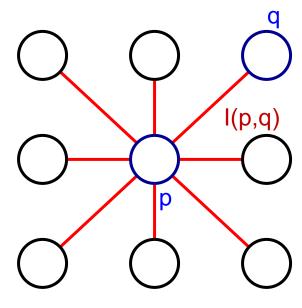
Intelligent Scissors

Treat the image as a graph

- Node for every pixel p
- Link between every adjacent pair of pixels, *p*,*q*
- Cost *I(p,q)* for each link

Cost combines several edge measures

- Laplacian Zero-Crossing, $f_z(q)$
- Gradient Magnitude, $f_G(q)$, Gradient Direction, $f_D(p,q)$ -



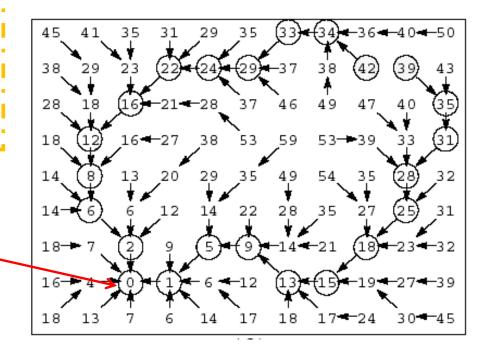
$$I(p,q) = \omega_Z \cdot f_Z(q) + \omega_D \cdot f_D(p,q) + \omega_G \cdot f_G(q)$$

p is an edge: we want q to be an edge too, with a high gradient, and the gradient direction to be similar at p and q



The LiveWire Algorithm

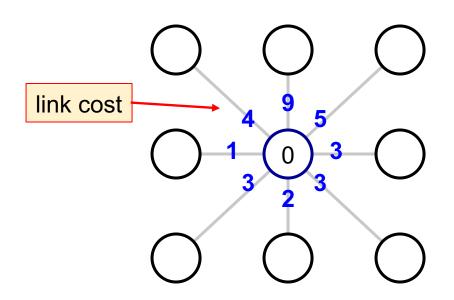
- Finds the minimum cost path between seed point and every point in the graph
- Output is a pointer at each pixel indicating the minimum cost path back to the seed
- As the mouse moves over the image the path from the current pixel is followed back to the seed



ACK: Prof. Tony Pridmore, UNUK

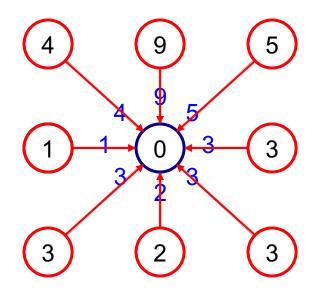
seed





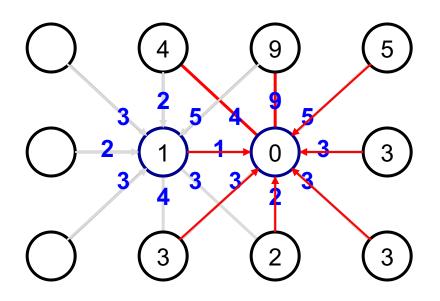
- 1. Init code cost to ∞ , set $p = seed\ point$, cost(p) = 0
- 2. Expand *p* as follows:
 - For each of p's neighbours q that are not expanded
 - ightharpoonup Set $cost(q) = min(cost(p) + c_{pq}, cost(q))$





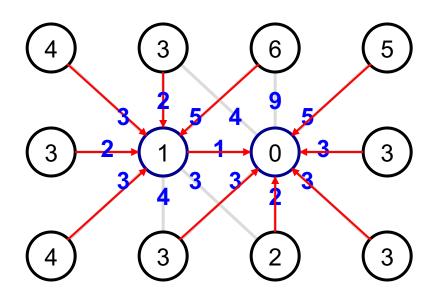
- 1. Init code cost to ∞ , set p = seed point, cost(p) = 0
- 2. Expand *p* as follows:
 - For each of p's neighbours q that are not expanded
 - ightharpoonup Set $cost(q) = min(cost(p) + c_{pq}, cost(q))$
 - √ if q's cost changed, make q point back to p
 - > Put q on the ACTIVE list (if not already there)





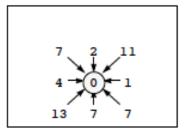
- 1. Init code cost to ∞ , set p = seed point, cost(p) = 0
- 2. Expand *p* as follows:
 - For each of p's neighbours q that are not expanded
 - ightharpoonup Set $cost(q) = min(cost(p) + c_{pq}, cost(q))$
 - √ if q's cost changed, make q point back to p
 - Put q on the ACTIVE list (if not already there)
- 3. Set r = node with minimum cost of the ACTIVE list
- 4. Repeat Step 2 for p = r

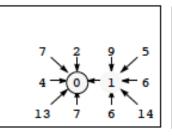


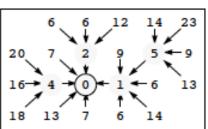


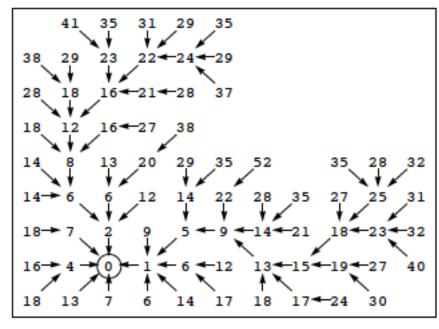
- 1. Init code cost to ∞ , set p = seed point, cost(p) = 0
- 2. Expand *p* as follows:
 - For each of p's neighbours q that are not expanded
 - ightharpoonup Set $cost(q) = min(cost(p) + c_{pq}, cost(q))$
 - ✓ if q's cost changed, make q point back to p
 - Put q on the ACTIVE list (if not already there)
- 3. Set r = node with minimum cost of the ACTIVE list
- 4. Repeat Step 2 for p = r

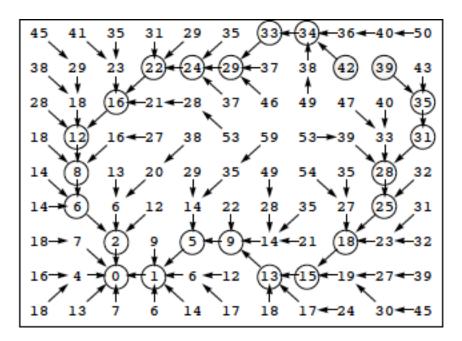








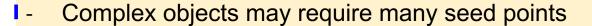






Cooling

The tool allows multiple seeds, but all seeds must be manually selected

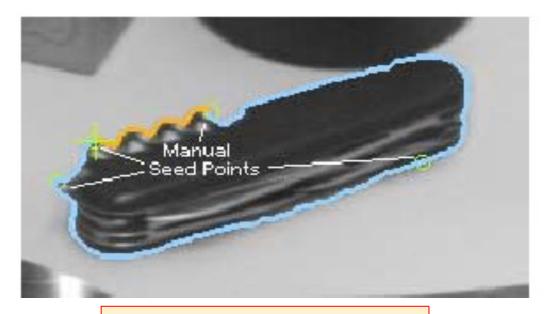




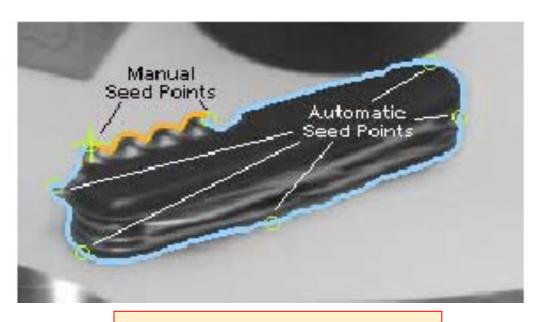
- As the user wraps the object, the early sections tend to (and need to) be fixed
- Path segments that remain fixed over N cursor movements are considered stable
- New seed points are created at the ends of stable segments



Cooling



Basic version: several manual seed points are needed to capture the shape



With cooling: two seeds are enough, the rest are created automatically



Interactive Dynamic Training

Some objects have stronger edges than other

If the desired edge is weaker than a nearby edge, then the path 'jumps' over to the stronger edge

Train the gradient magnitude to desire the weaker edge





- Use a sample of good path to train gradient magnitude
- i.e., change the weighting of the gradient magnitude term
- Update sample as path moves along the desired edge

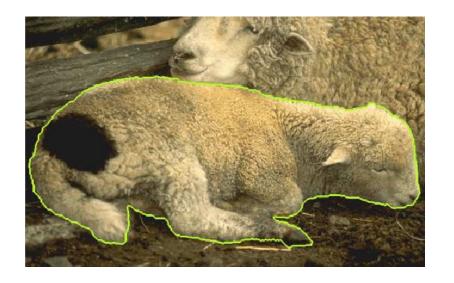
Allow user to enable and disable training as needed



Results







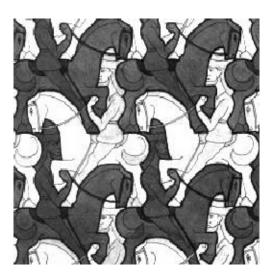


Blending Images



Blending



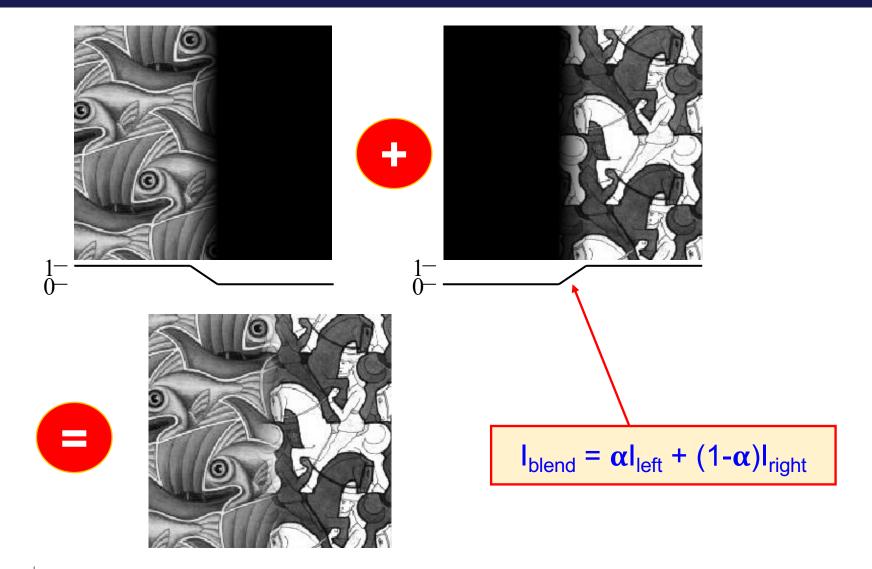


- Just putting sprites next to each other isn't enough
- Boundaries are clear, and attract the eye



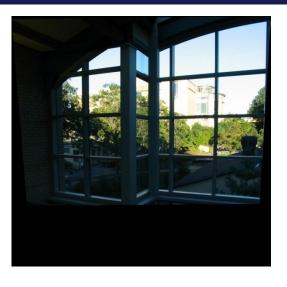


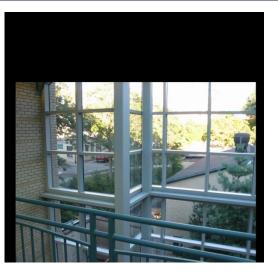
Alpha-Blending/Feathering

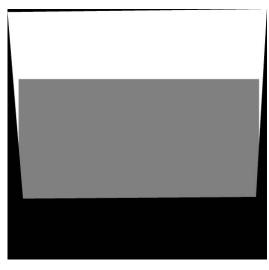




Setting α : Simple averaging







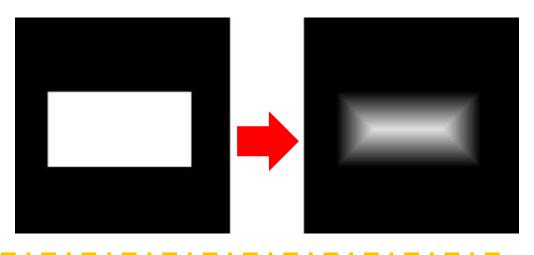


 α = 0.5 in the overlap region



Distance Transforms

Assign a non-negative integer each pixel giving the distance from that pixel to e.g., nearest black pixel



- A range of distance measure exist
- Characterise the shape of a binary image
- Can be used to guide blending

ACK: Prof. Tony Pridmore, UNUK

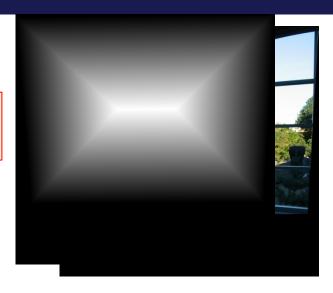
COMP 2032

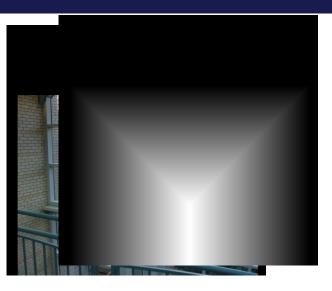
23

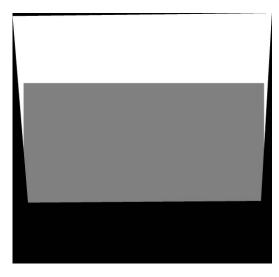


Setting α: Centre Seam

Distance Transform bwdist









 $\alpha = logical(dt1>dt2)$

Each pixel uses the image whose centre seam it is closest to

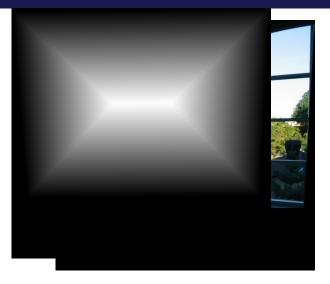
24

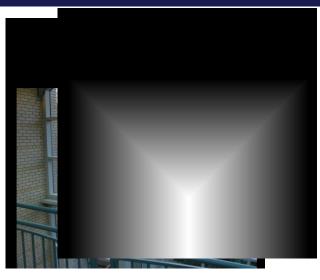
ACK: Prof. Tony Pridmore, UNUK

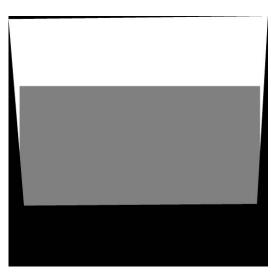
COMP 2032



Setting α: Centre Weighting

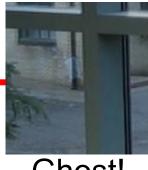








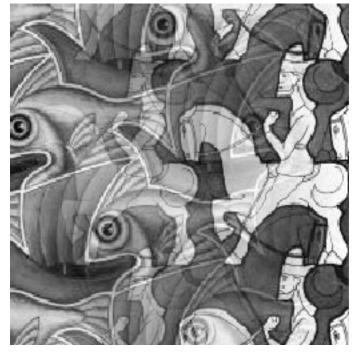
 $\alpha = dt1 / (dt1 + dt2)$

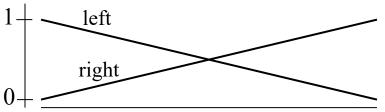


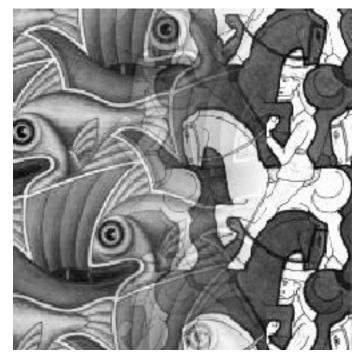
Ghost!

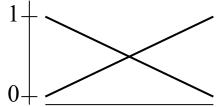


The Effect of Window Size





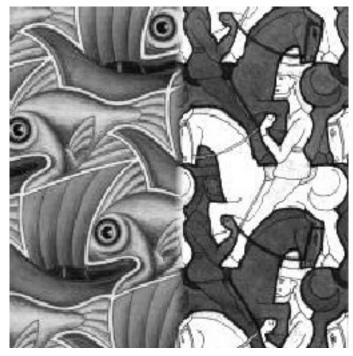




26



The Effect of Window Size





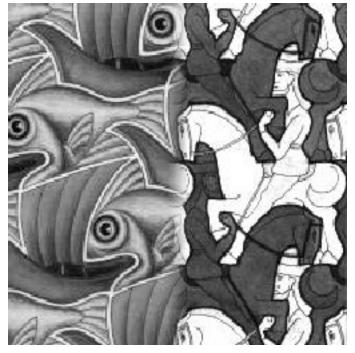




27



Good Window Size



 $0 + \underbrace{\bigvee_{0}}_{1}$

The optimal window size leaves the blend smooth, but not ghosted

To avoid seams:

Window = size of largest prominent feature

To avoid ghosting:

Window <= 2* size of largest prominent feature



Break





Geometric Transformations



Geometric Transformations

Transforms we have looked so far effect the content of the image array, but leave the spatial arrangement of the pixels intact, geometric or rubber-sheet transformations do not









Two stages:

Spatial transformation of pixel coordinates

1

Assign intensity values to new pixels

2



Spatial Transformations

Coordinate transforms are expressed as (x,y) = T(v,w)

e.g., (x,y) = (v/2, w/2) shrinks the image to half its size in both dimensions

One of the most common is the affine transform

$$[x \ y \ 1] = [v \ w \ 1] \ \mathbf{T} = [v \ w \ 1] \begin{bmatrix} t_{11} & t_{12} & 0 \\ t_{21} & t_{22} & 0 \\ t_{31} & t_{32} & 1 \end{bmatrix}$$

Which can scale, rotate, tanslate or shear an image

Those can be applied separately, or a sequence of transformations can be applied by a single 3 x 3 matrix

ACK: Prof. Tony Pridmore, UNUK

The order in which transforms are applied matters

Note



Transformation Name	Affine Matrix, T	Coordinate Equations	Example
Identity	$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$	x = v y = w	\int_{x}^{y}
Scaling	$\begin{bmatrix} c_x & 0 & 0 \\ 0 & c_y & 0 \\ 0 & 0 & 1 \end{bmatrix}$	$x = c_x v$ $y = c_y w$	
Rotation	$\begin{bmatrix} \cos \theta & \sin \theta & 0 \\ -\sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$	$x = v \cos \theta - w \sin \theta$ $y = v \cos \theta + w \sin \theta$	
Translation	$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ t_x & t_y & 1 \end{bmatrix}$	$x = v + t_x$ $y = w + t_y$	
Shear (vertical)	$\begin{bmatrix} 1 & 0 & 0 \\ s_v & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$	$x = v + s_v w$ $y = w$	
Shear (horizontal)	$\begin{bmatrix} 1 & s_h & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$	$x = v$ $y = s_h v + w$	



Applying Transformations

Forward Mapping

- Scan the input image, compute the (new) position of each pixel in the output image
- What do you do if two input pixels map to one output pixel?
- Some output locations may not be allocated a pixel at all

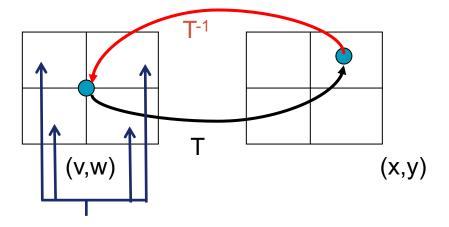
Inverse Mapping

- Scan the output image, use the inverse of the transformation
 (v,w) = T⁻¹(x,y) to work out which point in the input image maps to each output pixel location
- Use image interpolation to compute a new (output) pixel value from nearby (input) pixel values



Image Interpolation

In general, transformed pixel coordinates will not map neatly onto the coordinates of a regular rectangular array



We must use nearby (input) intensity values to compute a suitable output value at (x,y)



Image Interpolation

Nearest neighbour interpolation

Just pick the nearest pixel value and use that

Bilinear interpolation

take the four nearest neighbours and assume intensity varies linearly, then

$$v(x, y) = ax + by + cxy + d$$

Bicubic interpolation

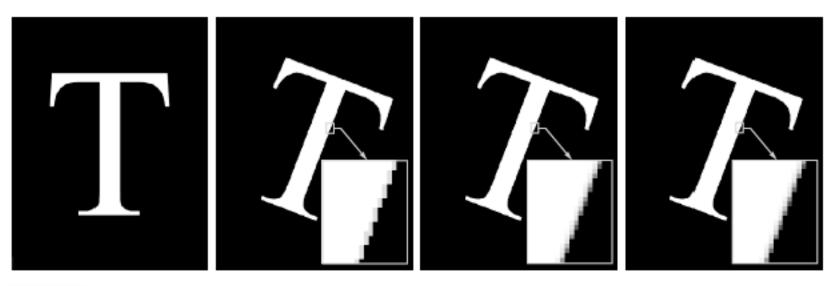
take the 16 nearest neighbours, fit a cubic, so

$$v(x, y) = \sum_{i=0}^{3} \sum_{j=0}^{3} a_{ij} x^{i} y^{j}$$

Bicubic interpolation is the standard, used in e.g., Photoshop



Image Interpolation



a b c d

FIGURE 2.36 (a) A 300 dpi image of the letter T. (b) Image rotated 21° using nearest neighbor interpolation to assign intensity values to the spatially transformed pixels. (c) Image rotated 21° using bilinear interpolation. (d) Image rotated 21° using bicubic interpolation. The enlarged sections show edge detail for the three interpolation approaches.

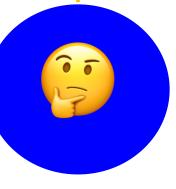
From Gonzalez & Woods

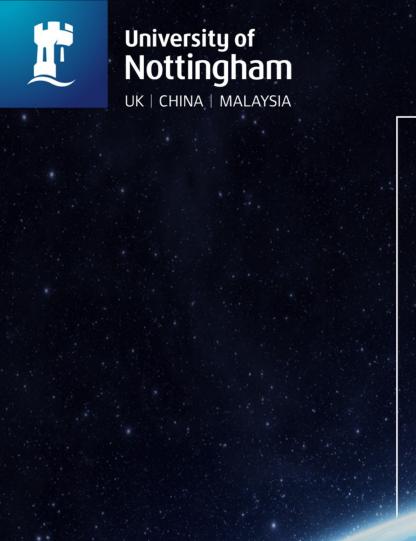


Summary



- 1. Livewire Algorithm
- 2. Blending Images
- 3. Geometric Transformation





Questions



NEXT:

Finale & Revision